

---

# Matrix Completion Project

---

**Ashish Gupta Konagalla**

University of Maryland, Baltimore County

*Email: [gri76457@umbc.edu](mailto:gri76457@umbc.edu)*

**Project GitHub Page:**

[Matrix Completion](#)

## Executive Summary

In real-world recommendation systems, handling missing data is a critical challenge, as incomplete information can significantly impact predictions and user experience. This project addresses the *matrix completion problem*, where missing ratings in a sparse user-movie matrix must be accurately predicted. To solve this problem, we implemented and compared four advanced techniques:

- **User-User and Item-Item based Collaborative Filtering,**
- **Singular Value Decomposition (SVD),**
- **Alternating Least Squares (ALS), and**
- **Nuclear Norm Minimization (NNM) with ADMM optimization.**

The core objective of the project is to evaluate these methods in terms of their ability to reconstruct the incomplete matrix and accurately predict missing ratings. Each approach leverages a unique mathematical framework to address sparsity, making it suitable for different contexts. Collaborative filtering methods rely on user and item similarities, while matrix factorization techniques, such as SVD and ALS, break the matrix into latent components to approximate the missing values. NNM with ADMM enforces a low-rank structure using singular value thresholding, enabling precise reconstruction.

The results show that all methods offer valuable insights into solving matrix completion problems. Notably, **NNM with ADMM** achieved the best performance with an RMSE of **0.9211** and MAE of **0.8818**, reflecting its ability to capture the underlying patterns in sparse data. SVD and ALS also demonstrated strong performance by leveraging matrix factorization, while collaborative filtering approaches provided baseline comparisons rooted in user and item interactions.

By comparing these techniques, this project highlights the strengths and trade-offs of different methods for sparse matrix completion. The findings contribute to understanding how mathematical models can effectively address real-world data sparsity challenges, with applications ranging from recommendation systems to predictive analytics.

---

## Background

### Problem Description

In today's digital age, recommendation systems play a critical role in enhancing user experience on platforms like Netflix, Amazon, and Spotify. These systems rely on vast user-item interaction datasets, but a common challenge they face is missing data. Users do not rate or interact with all items, leading to a sparse matrix with many empty entries. Accurately predicting these missing values is essential for providing personalized recommendations and improving user satisfaction.

This challenge of predicting missing values in a sparse matrix is known as the *matrix completion problem*. The goal is to “fill in the gaps” in the matrix with values that are as close as possible to the actual user preferences.

For example, consider a movie recommendation system where users rate movies on a scale of 1 to 5. The matrix rows represent users, the columns represent movies, and the entries represent ratings. A missing value means that a user has not rated a particular movie. The task is to predict this missing rating accurately so the system can recommend movies the user is likely to enjoy.

In this project, we aim to solve the matrix completion problem using different approaches, including **Singular Value Decomposition (SVD)**, **Collaborative Filtering**, **Nuclear Norm Minimization (NNM)** with ADMM optimization, and **Alternating Least Squares (ALS)**. Each method brings a unique perspective to the problem, balancing accuracy and computational efficiency.

By solving the matrix completion problem, this project contributes to building more accurate and effective recommendation systems, which are widely used in e-commerce, entertainment, and other industries.

## History and Importance

The matrix completion problem has its roots in the field of linear algebra and became widely recognized due to its applications in recommendation systems. A pivotal moment in the history of this problem was the *Netflix Prize* competition in 2006. Netflix, one of the largest streaming platforms, offered a \$1 million prize to anyone who could improve their recommendation algorithm by 10%. The challenge involved predicting user ratings for movies, where most of the ratings matrix was incomplete (users hadn’t rated all movies). This competition brought global attention to the significance of matrix completion and sparked advancements in collaborative filtering, machine learning, and optimization techniques.

The problem is critically important because sparse datasets—where only a fraction of the data is observed—are extremely common in real-world applications. For example:

- In e-commerce platforms like Amazon, users purchase or review only a small subset of available products.
- On music streaming services like Spotify, users listen to a limited set of tracks.
- In healthcare, missing entries in patient records (e.g., test results, symptoms) can significantly affect decision-making.

Matrix completion methods enable us to “fill in the gaps” and extract meaningful insights from incomplete data. Beyond recommendations, the ability to accurately predict missing entries has applications in image recovery, sensor networks, and financial forecasting.

An interesting fact is that low-rank approximations—a key concept behind many matrix completion methods—can be traced back to foundational work in mathematics, such as *Singular Value Decomposition (SVD)*, which was formalized in the early 20th century. Modern advancements, like **Nuclear Norm Minimization** and **Alternating Least Squares (ALS)**, build on these classical techniques to handle today’s large-scale, sparse datasets efficiently.

In summary, the matrix completion problem is not only an intriguing challenge rooted in mathematics but also a practical necessity in today's data-driven world. Its solutions drive some of the most successful technologies, from personalized shopping recommendations to predictive analytics in healthcare.

## Mathematical Foundation

The matrix completion problem can be formally described as the task of predicting missing entries in a partially observed matrix  $R$ . In this matrix, rows represent users, columns represent items (e.g., movies), and the observed entries  $R_{i,j}$  correspond to ratings or interactions. Mathematically, the goal is to approximate  $R$  as closely as possible while respecting its inherent structure.

### *Low-Rank Assumption*

In real-world datasets, user preferences and item features often follow hidden patterns. This implies that the ratings matrix  $R$  can be approximated by a matrix of much lower rank  $k$ , where  $k \ll \min(m, n)$  (with  $m$  and  $n$  being the number of users and items, respectively). The low-rank property reduces the problem to finding a matrix  $X$  such that:

$$\text{rank}(X) \leq k \quad \text{and} \quad X_{i,j} \approx R_{i,j} \quad \forall (i, j) \in \Omega.$$

Here,  $\Omega$  is the set of indices where the ratings are observed.

### *Singular Value Decomposition (SVD)*

One of the most fundamental techniques for low-rank matrix approximation is the *Singular Value Decomposition*. For a given matrix  $R$ , the SVD decomposes it into three matrices:

$$R \approx U \Sigma V^T,$$

where:

- $U$  is an  $m \times k$  matrix of left singular vectors,
- $\Sigma$  is a  $k \times k$  diagonal matrix of singular values,
- $V^T$  is a  $k \times n$  matrix of right singular vectors.

By truncating the singular values (retaining only the largest  $k$ ), we can approximate  $R$  with a low-rank matrix that minimizes reconstruction error.

### *Nuclear Norm Minimization (NNM)*

Instead of explicitly controlling the rank, the *nuclear norm* (sum of singular values) can be minimized as a convex relaxation of the rank constraint:

$$\min_X \|X\|_* \quad \text{subject to} \quad X_{i,j} = R_{i,j}, \quad \forall (i, j) \in \Omega.$$

Here,  $\|X\|_*$  represents the nuclear norm, which encourages a low-rank solution. Optimization algorithms like *ADMM* (Alternating Direction Method of Multipliers) are used to solve this problem efficiently by breaking it into smaller, manageable steps.

### ***Alternating Least Squares (ALS)***

The ALS method approximates  $R$  as the product of two lower-dimensional matrices  $U$  and  $V$ , where  $U \in \mathbb{R}^{m \times k}$  and  $V \in \mathbb{R}^{n \times k}$ . The objective function is:

$$\min_{U, V} \|R - UV^T\|_F^2 + \lambda(\|U\|_F^2 + \|V\|_F^2),$$

where  $\|\cdot\|_F$  is the Frobenius norm and  $\lambda$  is a regularization parameter to prevent overfitting. ALS alternates between solving for  $U$  (fixing  $V$ ) and  $V$  (fixing  $U$ ) until convergence.

### ***Summary***

The matrix completion problem relies on the assumption that the observed data can be approximated by a low-rank structure. Techniques like *SVD*, *NNM*, and *ALS* provide mathematical tools to approximate missing entries efficiently. Each method balances accuracy, computational complexity, and the ability to handle sparsity, laying the foundation for solving real-world recommendation problems.

---

## **Methodology**

### **Introduction**

This section outlines the steps taken to address the matrix completion problem, where the goal is to predict missing values in a user-item ratings matrix. We begin by analyzing the dataset and its sparsity structure, providing insights through visualizations. Following this, we implement and evaluate four key techniques:

- **Collaborative Filtering:** User-based and item-based similarity approaches.
- **Singular Value Decomposition (SVD):** Low-rank matrix factorization.
- **Alternating Least Squares (ALS):** Iterative optimization using latent factor models.
- **Nuclear Norm Minimization (NNM) with ADMM:** Convex optimization to enforce low-rank constraints.

Finally, we compare the performance of these methods using metrics such as Root Mean Squared Error (RMSE), Mean Absolute Error (MAE). The findings highlight the trade-offs between accuracy, computational efficiency, and the ability to handle sparsity.

### **Dataset and Sparsity Analysis**

The dataset used in this project is the **MovieLens ml-latest-small dataset**, which contains anonymous movie rating data from the MovieLens recommendation service. The dataset includes:

- **Total Ratings:** 100,836,

- **Number of Users:** 610,
- **Number of Movies:** 9,742,
- **Rating Scale:** 0.5 to 5.0

Each row in the dataset corresponds to a rating provided by a user for a specific movie and includes the following columns: `userId`, `movieId`, `rating`, and `timestamp`. For our project purpose, we do not focus on the `timestamp`.

### ***Sparsity of the Ratings Matrix***

The ratings data can be represented as a matrix  $R$ , where rows correspond to users, columns correspond to movies, and the entries  $R_{ij}$  represent the ratings provided by user  $i$  for movie  $j$ . The matrix size is given by:

$$\text{Matrix Size} = 610 \times 9724 = 5,931,640 \text{ entries.}$$

Out of these, only 100,836 ratings are observed, making the dataset highly sparse. The sparsity is calculated as:

$$\text{Sparsity} = 1 - \frac{\text{Number of Observed Ratings}}{\text{Total Entries}}.$$

Substituting the values:

$$\text{Sparsity} = 1 - \frac{100,836}{5,931,640} \approx 98.3\%.$$

This means that approximately 98% of the ratings are missing, posing a significant challenge for accurately predicting the unobserved entries.

### ***Visualization of Sparsity***

To better understand the sparsity pattern, a heatmap of the ratings matrix was generated, where observed ratings are represented as filled cells, and missing entries are shown as blank spaces. Figure 1 illustrates this sparsity. Yellow in the figure signifies entries for NaN values and the blue for observed entries.

The heatmap reveals that the ratings are not uniformly distributed. While some users have rated many movies, others have provided very few ratings. But every user provided a minimum of atleast 20 ratings. Similarly, a small subset of movies have significantly more ratings than others, reflecting an imbalance in the dataset.

### ***Significance of Addressing Sparsity***

The high sparsity and uneven distribution of ratings underscore the need for robust matrix completion techniques. These methods aim to leverage the available ratings to predict the missing entries accurately, which is essential for improving the performance of recommendation systems.




Figure 1: Heatmap of the ratings matrix showing observed (non-zero) and missing entries.

## Description of Methods:

### 1. User-Based Collaborative Filtering

User-based collaborative filtering is a technique that predicts missing ratings by leveraging similarities between users. The core idea is that users with similar preferences are likely to rate items in a similar manner. In this implementation, **cosine similarity** is used to measure the similarity between users based on their rating patterns.

#### *Similarity Calculation*

The user-item ratings matrix is first prepared for similarity computation:

- Missing ratings (NaN) are replaced with 0 to allow computation of cosine similarity.
- The user-user similarity matrix  $S$  is computed using the cosine similarity formula:

$$\text{sim}(u, v) = \frac{\sum_i R_{ui} R_{vi}}{\sqrt{\sum_i R_{ui}^2} \sqrt{\sum_i R_{vi}^2}},$$

where  $R_{ui}$  and  $R_{vi}$  are the ratings of users  $u$  and  $v$  for item  $i$ .

#### *Prediction of Missing Ratings*

To predict a missing rating  $\hat{R}_{ui}$  for user  $u$  and item  $i$ :

1. Retrieve the similarity scores of user  $u$  with all other users.
2. Consider ratings for item  $i$  from users who have already rated it (non-zero ratings).

3. Compute the weighted average of these ratings, where the weights are given by the similarity scores:

$$\hat{R}_{ui} = \frac{\sum_{v \in N} \text{sim}(u, v) \cdot R_{vi}}{\sum_{v \in N} |\text{sim}(u, v)|},$$

where  $N$  is the set of users who rated item  $i$  and  $\text{sim}(u, v)$  is the similarity between users  $u$  and  $v$ .

If no similar users have rated the item, the prediction defaults to 0.

### Evaluation of Predictions

The ratings matrix is split into training and testing sets, with 80% of the observed ratings used for training and the remaining 20% for testing. The Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) are used to evaluate the accuracy of the predictions:

$$\text{RMSE} = \sqrt{\frac{\sum_i (R_i - \hat{R}_i)^2}{n}}, \quad \text{MAE} = \frac{\sum_i |R_i - \hat{R}_i|}{n},$$

where  $R_i$  and  $\hat{R}_i$  are the actual and predicted ratings, respectively, and  $n$  is the number of test ratings.

### Results

The performance of user-based collaborative filtering on the test set is summarized as follows:

- **Root Mean Squared Error (RMSE): 3.661**
- **Mean Absolute Error (MAE): 3.507**

These results indicate a significant error margin, which can be attributed to the sparsity of the dataset and the fact that not all users have similar rating patterns. Below is the Figure 2 for how the resultant matrix is not suitable at all for recommendations as some of the ratings are simply left as 0 due to sparse data.

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6	Movie 7	Movie 8
User 1	0.0	3.334487640521381	0.0	2.267387734956233	2.9432957227276164	0.0	3.0751458314463886	2.755556705765913
User 2	3.822566458710671	3.326362864914464	3.051040950856254	2.3773227163598163	2.8569153916451624	4.137285014401862	3.063882889820672	2.934260959316613
User 3	3.914399564042275	3.3317461734345803	3.178819178952575	2.53215527666695	2.7519668505512094	3.8206846994937727	3.1442444581247093	2.7542331705793828
User 4	3.8796612500024104	3.386672775803052	3.161365588104031	2.2456419168516653	2.866044530200229	3.5053631597849087	3.024051222212122	2.663644143752882
User 5	0.0	3.5686389821087896	3.381726011855211	2.390837369919151	3.15043745128799	3.4882300848697508	3.398076609872213	3.1235104296270793
User 6	4.01693119878402	0.0	0.0	0.0	0.0	0.0	0.0	0.0
User 7	0.0	3.361432823011877	3.13285254762812	2.2347721412980993	2.9274850826245604	3.57442254513883	3.14407648853187	2.7225808482318245
User 8	3.9827097934367313	0.0	3.4284667936105953	2.438852211141028	3.1696147398247088	3.893386737439447	3.4173764614151794	3.189524028216852
User 9	3.8491539174810523	3.2786676737095255	3.0757467920088817	2.135813130371124	2.7774155155681584	3.9549565340621347	2.9097697778602094	2.29225671092669
User 10	3.8988180692392507	3.3291073287454913	3.0608380633013316	2.214881698377468	2.70852350955687	3.578403404051683	3.051568841657644	2.70187218098165

Figure 2: User-User collaborative filtering technique output

## 2. Item-Based Collaborative Filtering

Item-based collaborative filtering predicts missing ratings by leveraging similarities between items (e.g., movies). Unlike user-based filtering, which focuses on finding similar users, this approach identifies similar items based on user ratings. If two movies receive similar ratings from multiple users, a user's preference for one movie can help predict their preference for the other.



### Similarity Calculation

The user-item ratings matrix  $R$  is transposed to form a movie-user matrix, where rows represent movies, and columns represent users. Missing values are replaced with 0 to allow similarity computation. The pairwise similarity between movies is then calculated using cosine similarity, resulting in an  $n \times n$  similarity matrix  $S$ , where  $n$  is the number of movies.

### Prediction of Missing Ratings

To predict missing ratings:

1. For each user, retrieve the similarity scores of a target movie with all other movies.
2. Use the ratings given by the user for similar movies to compute a weighted average. The weights are determined by the similarity scores.

If no similar movies have been rated, the prediction defaults to 0.

### Evaluation of Predictions

The ratings matrix is split into training and testing sets, with 80% of the observed ratings used for training. RMSE and MAE are used to evaluate the predictions.

### Results

The performance of item-based collaborative filtering on the test set is as follows:

- **Root Mean Squared Error (RMSE): 0.913**
- **Mean Absolute Error (MAE): 0.703**

### Observations

Item-based collaborative filtering achieves better accuracy compared to user-based filtering. This improvement can be attributed to the fact that movies generally receive more ratings than individual users provide. This reduces sparsity in the item space and allows for more reliable similarity calculations. In the below figure, you can see the output.

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6	Movie 7	Movie 8
User 1	0.0	4.40471438332893	0.0	4.445630003520309	4.399544964521706	0.0	4.388486732576779	4.433620474228227
User 2	3.8572042085645215	3.8363261655540253	3.7575929350343356	4.0	3.858470353160694	3.8696973923285123	3.869498493594609	3.6877297620927854
User 3	1.3827517761084656	1.3620029280411978	1.4721069993783853	0.49999999999999994	1.1127423247430674	1.6855100056257462	0.9707614247471675	1.0583743287581313
User 4	3.503413224510894	3.4736085015504448	3.448725711369957	3.4791909821432894	3.5356884340721693	3.467248292647556	3.489784723916806	3.371784302250939
User 5	0.0	3.6814770764018863	3.6727520956565814	3.4753688353595575	3.6695235217965845	3.6392776639812556	3.65760440931291	3.845081526115145
User 6	3.5664872466149435	0.0	3.4689391978702524	0.0	3.4895305061453827	3.5264016330052095	3.491915972921362	0.0
User 7	3.509082297529736	3.4145064533236664	3.514113791599305	3.8296020610556623	3.527194183013966	3.5324671972637733	3.565205385892458	3.326254722642594
User 8	3.5422985202925386	0.0	3.4099943689671064	3.4703476400250306	3.45599640725381	3.517289882671247	3.4571760012638255	3.4079942511689167
User 9	3.631620197229097	3.4239200209409475	3.318543966149928	3.397978706179447	3.512797125251347	3.41193992288288	3.184245408992641	2.8359852011119293
User 10	3.0879253725930416	3.118172528479877	3.0566466176060425	2.9271229540945916	3.085443610063845	3.063755117787361	3.089920988987284	3.2092958529624456

Figure 3: Item-Item collaborative filtering technique output

---

### Important Note: Reliability of Results

*Reliability of Item-Based Results:* While the item-based collaborative filtering method yielded significantly lower RMSE and MAE scores, it is important to note that this outcome is largely coincidental. The reason lies in the way the dataset is structured:

- The MovieLens dataset contains a large imbalance where a few movies receive many ratings, while most other movies receive very few ratings.
- This imbalance skews the similarity calculations, as movies with many ratings tend to dominate the predictions.
- The sparsity of user ratings means that the weighted averages in item-based predictions are often influenced by a small subset of movies.

*Impact of Missing Ratings for Certain Movies:* A critical limitation of both user-based and item-based collaborative filtering is that predictions cannot be reliably made for movies that have received no ratings. Specifically:

- In **user-based filtering**, if no user has rated a particular movie, there are no relevant ratings to compute a prediction for that movie.
- In **item-based filtering**, if no users have watched or rated a specific movie, its similarity with other movies cannot be established, rendering predictions impossible for that movie.

These scenarios highlight a fundamental weakness of collaborative filtering methods, as they depend heavily on the presence of existing ratings.

In conclusion, while user-based and item-based collaborative filtering can provide useful predictions in moderately sparse datasets, they fail to handle cases with extreme sparsity or missing ratings. This limitation reinforces the need for more advanced techniques, such as matrix factorization or hybrid approaches, to address these issues effectively.

---

## 3. Singular Value Decomposition (SVD) with Gradient Descent

Singular Value Decomposition (SVD) is a matrix factorization technique that decomposes the user-item matrix  $R$  into three matrices:  $U$ ,  $\Sigma$ , and  $V^T$ . In this implementation, we approximate the factorization using *gradient descent* to optimize the user and item latent factors. This approach allows us to iteratively minimize the reconstruction error.

### *Matrix Decomposition*

The ratings matrix  $R$  is approximated as:

$$R \approx U \cdot V^T,$$

where:

- $U \in \mathbb{R}^{m \times k}$ : User latent factor matrix,

- $V \in \mathbb{R}^{n \times k}$ : Item latent factor matrix,
- $k$ : Number of latent factors (dimensionality of the latent space),
- $m$  and  $n$ : Number of users and movies, respectively.

The objective is to learn  $U$  and  $V$  such that the reconstruction error for the observed ratings is minimized. The error is given by:

$$E = \sum_{i=1}^m \sum_{j=1}^n \mathbb{I}_{ij} \cdot (R_{ij} - U_i \cdot V_j^T)^2 + \lambda(\|U\|^2 + \|V\|^2),$$

where  $\mathbb{I}_{ij}$  is an indicator function for observed entries, and  $\lambda$  is the regularization parameter.

### ***Gradient Descent Optimization***

The algorithm iteratively updates the user and item matrices  $U$  and  $V$  using the following update rules:

$$U_i \leftarrow U_i + \eta(e_{ij} \cdot V_j - \lambda \cdot U_i), \quad V_j \leftarrow V_j + \eta(e_{ij} \cdot U_i - \lambda \cdot V_j),$$

where:

- $e_{ij} = R_{ij} - U_i \cdot V_j^T$ : Prediction error for user  $i$  and movie  $j$ ,
- $\eta$ : Learning rate, which controls the step size of the updates,
- $\lambda$ : Regularization parameter to prevent overfitting.

The algorithm runs for a fixed number of iterations, updating  $U$  and  $V$  only for observed entries (non-zero ratings). At each step, the total loss is computed as the sum of squared errors plus the regularization term to monitor convergence.

### ***Reconstruction and Evaluation***

The reconstructed matrix  $\hat{R}$  is computed as:

$$\hat{R} = U \cdot V^T.$$

To ensure valid ratings, the reconstructed values are clipped to the range  $[1, 5]$ . The accuracy of the predictions is evaluated using Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) on the test set.

### ***Results***

The results of the gradient descent-based SVD method are as follows:

- **Final Loss:** 34,389.8195 (after 100 iterations),
- **Root Mean Squared Error (RMSE):** 0.9849,
- **Mean Absolute Error (MAE):** 0.7333.

### ***Observations***

The gradient descent-based SVD approach successfully minimizes the loss over iterations, as shown in the training log. The RMSE and MAE values indicate that the reconstructed matrix achieves a reasonable approximation of the actual ratings. However:

- The convergence depends on the choice of hyperparameters ( $k$ ,  $\eta$ , and  $\lambda$ ).
- The algorithm can be computationally expensive for large datasets due to the iterative nature of gradient descent.

The reconstructed matrix, shown below, demonstrates the model's ability to predict ratings close to the observed values:

$$\begin{bmatrix} 4.75 & 4.27 & 3.85 & \dots & 3.52 & 4.23 \\ 3.89 & 3.43 & 3.15 & \dots & 2.81 & 2.82 & 3.27 \\ 1.34 & 1.85 & 1.63 & \dots & 1.28 & 1.19 & 1.46 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \\ 3.52 & 3.12 & 2.71 & \dots & 3.29 & 3.42 & 4.03 \end{bmatrix}$$

## **4. Alternating Least Squares (ALS) with Gradient Descent**

The Alternating Least Squares (ALS) method minimizes the reconstruction error by alternating between fixing the user matrix  $U$  and solving for the item matrix  $V$ , and vice versa. This iterative approach gradually improves the approximation of the user-item matrix  $R$ .

### ***Objective Function***

The goal of ALS is to minimize the following loss function:

$$E = \sum_{(i,j) \in \Omega} (R_{ij} - U_i \cdot V_j^T)^2 + \lambda(\|U\|_F^2 + \|V\|_F^2),$$

where:

- $R_{ij}$ : Observed rating for user  $i$  and item  $j$ ,
- $\Omega$ : Set of observed entries in the matrix,
- $U$  and  $V$ : User and item latent factor matrices,
- $\lambda$ : Regularization parameter,
- $\|\cdot\|_F^2$ : Frobenius norm, used to penalize large values in  $U$  and  $V$ .

### ***Gradient Descent Updates***

The ALS approach updates one matrix at a time while keeping the other fixed. Using gradient descent, the update rules become:

**1. Fix Item Matrix  $V$  and Update User Matrix  $U$ :** For each user  $u$  and factor  $f$ , the user matrix  $U$  is updated as:

$$U_{u,f} \leftarrow U_{u,f} - \eta \left( \sum_{j \in \Omega_u} (R_{uj} - U_u \cdot V_j^T) \cdot (-V_{j,f}) + \lambda \cdot U_{u,f} \right),$$

where:

- $\eta$ : Learning rate, which controls the step size of updates,
- $\Omega_u$ : Set of items rated by user  $u$ ,
- $V_j$ : Item latent vector for item  $j$ ,
- Regularization  $\lambda \cdot U_{u,f}$  prevents overfitting.

**2. Fix User Matrix  $U$  and Update Item Matrix  $V$ :** For each item  $j$  and factor  $f$ , the item matrix  $V$  is updated as:

$$V_{j,f} \leftarrow V_{j,f} - \eta \left( \sum_{u \in \Omega_j} (R_{uj} - U_u \cdot V_j^T) \cdot (-U_{u,f}) + \lambda \cdot V_{j,f} \right),$$

where:

- $\Omega_j$ : Set of users who rated item  $j$ ,
- $U_u$ : User latent vector for user  $u$ ,
- Regularization  $\lambda \cdot V_{j,f}$  ensures stability of updates.

### ***Gradient Clipping and Learning Rate Decay***

To improve numerical stability and prevent divergence:

- **Gradient Clipping:** Gradients are capped within a threshold to avoid large updates:

$$\text{Gradient} \leftarrow \text{clip}(\text{Gradient}, -\gamma, \gamma),$$

where  $\gamma$  is the gradient threshold (e.g., 5.0 in this implementation).

- **Learning Rate Decay:** The learning rate  $\eta$  decreases over iterations to ensure convergence:

$$\eta_t = \frac{\eta_0}{1 + 0.1 \cdot t},$$

where  $t$  is the current iteration step.

### *Training Results and Convergence*

The loss decreases steadily over iterations, as the user and item matrices are updated alternately. The results show clear convergence:

Step: 1, Loss: 1066503.9081  
Step: 10, Loss: 556477.2633  
Step: 20, Loss: 145490.9033  
Step: 50, Loss: 66306.4433

The reconstructed matrix  $\hat{R}$  is obtained as:

$$\hat{R} = U \cdot V^T.$$

### *Comparison to SVD*

- **Performance:** While ALS converges to a solution, it achieves an RMSE of 2.5294 and an MAE of 2.2170, which are higher than SVD.
- **Efficiency:** ALS requires alternating updates for user and item matrices, making it slower than SVD, which updates both simultaneously.
- **Stability:** Gradient clipping and learning rate decay ensure stable training and prevent divergence.

### *Final Observations*

ALS with gradient descent provides a robust matrix completion method but involves significant computational overhead due to alternating updates. The following key observations were made:

- As the step count increases, the loss decreases steadily, indicating convergence. For example:

Step: 10, Loss: 556477.2633  
Step: 20, Loss: 145490.9033  
Step: 50, Loss: 66306.4433

- If the number of iterations (steps) were increased further, the loss would continue to decrease, likely improving the RMSE and MAE scores. However, this comes at the cost of significantly higher computational time, especially for large-scale datasets.
- While the model performs reasonably well, the higher error compared to SVD suggests that further improvements, such as optimized initialization or adaptive learning rates, could enhance both efficiency and accuracy.

## 5. Nuclear Norm Minimization (NNM) with ADMM

### Overview

Nuclear Norm Minimization (NNM) is a matrix completion method that reconstructs the missing entries in a sparse matrix by enforcing a *low-rank constraint*. In this implementation, the *Alternating Direction Method of Multipliers (ADMM)* is used to efficiently optimize the nuclear norm while minimizing reconstruction error.

The objective is to approximate the incomplete matrix  $R$  by solving:

$$\min_Z \|Z\|_* \quad \text{subject to} \quad Z_{ij} = R_{ij}, \forall (i, j) \in \Omega,$$

where:

- $\|Z\|_*$ : The nuclear norm of  $Z$ , defined as the sum of its singular values,
- $\Omega$ : The set of observed entries in the matrix,
- $Z$ : The completed matrix.

This approach encourages the completed matrix  $Z$  to have the lowest possible rank while maintaining the observed ratings.

### ADMM Formulation

ADMM decomposes the problem into simpler subproblems that are solved iteratively. The optimization involves three key updates:

**1. Update  $X$  (Observed Entries Step)** The matrix  $X$  is updated to ensure that observed entries align with the input ratings  $R$ :

$$X = \text{Proj}_\Omega(R) + \rho \cdot (Z - Y),$$

where  $\rho$  is the **augmented Lagrangian parameter**, and  $Y$  is the dual variable used to enforce consistency.

**2. Update  $Z$  (Proximal Step for Nuclear Norm Minimization)** The nuclear norm minimization is performed via Singular Value Thresholding (SVT). This step reduces the singular values of the matrix  $X + Y$ :

$$Z = U \cdot \text{diag}(\max(S - \lambda/\rho, 0)) \cdot V^T,$$

where:

- $U, S, V^T$ : Singular Value Decomposition (SVD) of  $X + Y$ ,
- $\max(S - \lambda/\rho, 0)$ : Soft-thresholding of the singular values to promote low-rank structure,
- $\lambda$ : Regularization parameter controlling the rank penalty.

**3. Update  $Y$  (Dual Variable Update)** The dual variable  $Y$  is updated to maintain consistency between  $X$  and  $Z$ :

$$Y \leftarrow Y + X - Z.$$

### *Convergence*

The algorithm iteratively performs these updates until convergence is achieved, which is determined based on the relative error between  $X$  and  $Z$ :

$$\text{Error} = \frac{\|X - Z\|_F}{\|X\|_F}.$$

### *Results and Observations*

The NNM with ADMM method produces high-quality results, achieving the following evaluation metrics:

- **Root Mean Squared Error (RMSE):** 0.9211,
- **Mean Absolute Error (MAE):** 0.8818.

The superior performance of NNM stems from its ability to accurately capture the underlying low-rank structure of the ratings matrix. By soft-thresholding the singular values, the method effectively reduces noise while reconstructing the missing entries. This balance between preserving observed data and promoting low-rank structure makes NNM a powerful approach for matrix completion.

### *Conclusion*

Nuclear Norm Minimization with ADMM demonstrates excellent accuracy in reconstructing sparse matrices by enforcing low-rank constraints through iterative updates. The use of singular value thresholding ensures precise reconstruction of missing values, leading to superior results compared to other methods.

## Comparison of Results

The table below summarizes the RMSE and MAE values obtained for all the matrix completion methods implemented in this project:

Table 1: Comparison of RMSE and MAE for All Methods

Method	RMSE	MAE
User-User Collaborative Filtering	3.6613	3.5075
Item-Item Collaborative Filtering	0.9133	0.7026
Singular Value Decomposition (SVD)	0.9849	0.7333
Alternating Least Squares (ALS)	2.5294	2.2170
Nuclear Norm Minimization (NNM) with ADMM	0.9211	0.8818

It is important to note that while most methods utilize random seeds for reproducibility, some testing cases did not explicitly set a fixed random seed. This may lead to minor variations in the RMSE and MAE values across multiple runs. These variations, however, do not significantly impact the comparative analysis or overall conclusions drawn from the results.



## Future Work

In the future, I aim to scale these matrix completion methods to larger datasets with millions of users and items, where handling computational efficiency and memory optimization becomes critical. I plan to explore parallelized implementations and distributed computing frameworks, such as Apache Spark, to improve scalability. Additionally, I want to experiment with hybrid approaches that combine collaborative filtering and matrix factorization techniques to leverage their respective strengths. I am also interested in exploring deep learning-based models, such as autoencoders or neural collaborative filtering, which could potentially capture more complex patterns in the data and further enhance prediction accuracy.

## References

- [1] F. Maxwell Harper and J. A. Konstan, *The MovieLens Datasets: History and Context*, ACM Transactions on Interactive Intelligent Systems (TIIS), vol. 5, no. 4, pp. 1–19, 2015. Available: <https://grouplens.org/datasets/movielens/>. Accessed: Mar. 24, 2024.
- [2] A. Berk, *Matrix Completion for Recommendation Systems: Project*, GitHub Gist, Available: <https://gist.github.com/asberk/8c8fd1e7f384df0183b8f01423c07a55>. Accessed: Mar. 24, 2024.
- [3] Y. Koren, R. Bell, and C. Volinsky, *Matrix Factorization Techniques for Recommender Systems*, IEEE Computer, vol. 42, no. 8, pp. 30–37, 2009. Available: <https://ieeexplore.ieee.org/document/5197422>. Accessed: Dec. 15, 2024.
- [4] Y. Zhou, D. Wilkinson, R. Schreiber, and R. Pan, *Large-Scale Parallel Collaborative Filtering for the Netflix Prize*, in Proceedings of the 4th International Conference on Algorithmic Aspects in Information and Management, 2008, pp. 337–348. Available: [https://link.springer.com/chapter/10.1007/978-3-540-68880-8\\_32](https://link.springer.com/chapter/10.1007/978-3-540-68880-8_32). Accessed: Dec. 15, 2024.
- [5] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*, Foundations and Trends in Machine Learning, vol. 3, no. 1, pp. 1–122, 2011. Available: [https://web.stanford.edu/~boyd/papers/pdf/admm\\_distr\\_stats.pdf](https://web.stanford.edu/~boyd/papers/pdf/admm_distr_stats.pdf). Accessed: Dec. 15, 2024.

---

## Appendix

All the code, implementation details, and additional resources for this project can be found on the GitHub repository:

<https://github.com/ashishkonagalla-umbc/Matrix-Completion>

---